


GPU Speed Of Light

All

High-level overview of the utilization for compute and memory resources of the GPU. For each unit, the Speed Of Light (SOL) reports the achieved percentage of utilization with respect to the theoretical maximum.

SOL SM [%]	84.09	Duration [nsecond]	732352
SOL Memory [%]	6.53	Elapsed Cycles [cycle]	428833
SOL TEX [%]	5.83	SM Active Cycles [cycle]	421051.35
SOL L2 [%]	6.53	SM Frequency [cycle/second]	585550117.98
SOL FB [%]	6.04	Memory Frequency [cycle/second]	1248623612.69



SOL SM Breakdown		SOL Memory Breakdown	
SOL SM: Pipe Fp64 Cycles Active [%]	84.09	SOL L2: T Sectors [%]	6.53
SOL SM: Mio Pq Write Cycles Active [%]	26.28	SOL GPU: Dram Throughput [%]	6.04
SOL SM: Mio Inst Issued [%]	5.26	SOL L1: M Xbar2l1tex Read Sectors [%]	5.72
SOL SM: Inst Executed Pipe Lsu [%]	4.30	SOL L1: Data Pipe Lsu Wavefronts [%]	4.93
SOL SM: Issue Active [%]	3.89	SOL L2: Lts2xbar Cycles Active [%]	4.90
SOL SM: Inst Executed [%]	3.88	SOL L1: Lsu Writeback Active [%]	4.45
SOL SM: Mio2rf Writeback Active [%]	2.03	SOL L1: Lsuin Requests [%]	4.30
SOL SM: Pipe Fma Cycles Active [%]	2.03	SOL L1: M L1tex2xbar Req Cycles Active [%]	2.64
SOL SM: Inst Executed Pipe Cbu Pred On Any [%]	0.97	SOL L2: Xbar2lts Cycles Active [%]	2.26
SOL SM: Inst Executed Pipe Xu [%]	0.96	SOL L2: D Sectors [%]	2.12
SOL SM: Inst Executed Pipe Adu [%]	0.96	SOL L2: T Tag Requests [%]	1.64
SOL SM: Pipe Alu Cycles Active [%]	0.72	SOL L1: Data Bank Reads [%]	1.35
SOL SM: Mio Pq Read Cycles Active [%]	0.48	SOL L1: Data Bank Writes [%]	1.27
SOL SM: Inst Executed Pipe lpa [%]	0	SOL L2: D Sectors Fill Device [%]	0.25
SOL SM: Inst Executed Pipe Tex [%]	0	SOL L1: Texin Sm2tex Req Cycles Active [%]	0.00
SOL SM: Inst Executed Pipe Uniform [%]	0	SOL L1: F Wavefronts [%]	0.00
SOL SM: Inst Executed Pipe Fp16 [%]	0	SOL L2: D Atomic Input Cycles Active [%]	0
SOL IDC: Request Cycles Active [%]	0	SOL L2: D Sectors Fill Systemem [%]	0
SOL SM: Pipe Shared Cycles Active [%]	0	SOL L1: Tex Writeback Active [%]	0
SOL SM: Pipe Tensor Cycles Active [%]	0	SOL L1: Data Pipe Tex Wavefronts [%]	0

Launch Statistics

Summary of the configuration used to launch the kernel. The launch configuration defines the size of the kernel grid, the division of the grid into blocks, and the GPU resources needed to execute the kernel. Choosing an efficient launch configuration maximizes device utilization.

Grid Size	1280	Registers Per Thread [register/thread]	30
Block Size	1024	Static Shared Memory Per Block [byte/block]	0
Threads [thread]	1310720	Dynamic Shared Memory Per Block [byte/block]	0
Waves Per SM	32	Shared Memory Configuration Size [byte]	32768

Occupancy

Occupancy is the ratio of the number of active warps per multiprocessor to the maximum number of possible active warps. Another way to view occupancy is the percentage of the hardware's ability to process warps that is actively in use. Higher occupancy does not always result in higher performance, however, low occupancy always reduces the ability to hide latencies, resulting in overall performance degradation. Large discrepancies between the theoretical and the achieved occupancy during execution typically indicates highly imbalanced workloads.

Theoretical Occupancy [%]	100	Block Limit Registers [block]	2
Theoretical Active Warps per SM [warp/cycle]	32	Block Limit Shared Mem [block]	16
Achieved Occupancy [%]	94.57	Block Limit Warps [block]	1
Achieved Active Warps Per SM [warp]	30.26	Block Limit SM [block]	16